


Structured P2P: DHT Approaches

- ❑ DHT service and issues
- ❑ CARP
- ❑ Consistent Hashing
- ❑ Chord
- ❑ CAN
- ❑ Pastry/Tapestry
- ❑ Hierarchical lookup services

Notes based on notes by K.W. Ross, J. Kurose, D. Rubenstein, and others

Challenge: Locating Content



- ❑ Simplest strategy: expanding ring search
- ❑ If K of N nodes have copy, expected search cost *at least* N/K , i.e., $O(N)$
- ❑ Need many cached copies to keep search overhead small

Directed Searches

❑ Idea:

- assign particular nodes to hold particular content (or pointers to it, like an information booth)
- when a node wants that content, go to the node that is supposed to have or know about it

❑ Challenges:


- Distributed: want to distribute responsibilities among existing nodes in the overlay
- Adaptive: nodes join and leave the P2P overlay
 - distribute knowledge responsibility to joining nodes
 - redistribute responsibility knowledge from leaving nodes

Distributed Hash Table (DHT)

- ❑ DHT = distributed P2P database
- ❑ Database has (key, value) pairs;
 - key: ss number; value: human name
 - key: content type; value: IP address
- ❑ Peers query DB with key
 - DB returns values that match the key
- ❑ Peers can also insert (key, value) peers

DHT Step 1: The Hash

- Introduce a hash function to map the object being searched for to a unique identifier:
 - e.g., $h(\text{"NGC'02 Tutorial Notes"}) \rightarrow 8045$
- Distribute the range of the hash function among all nodes in the network




- Each node must "know about" at least one copy of each object that hashes within its range (when one exists)

"Knowing about objects"

□ Two alternatives

- Node can cache each (existing) object that hashes within its range
- Pointer-based: level of indirection - node caches pointer to location(s) of object



DHT Step 2: Routing

- ❑ For each object, node(s) whose range(s) cover that object must be reachable via a “short” path
 - by the querier node (assumed can be chosen arbitrarily)
 - by nodes that have copies of the object (when pointer-based approach is used)
- ❑ The different approaches (CAN, Chord, Pastry, Tapestry) differ fundamentally only in the routing approach
 - any “good” random hash function will suffice

DHT Routing: Other Challenges


- ❑ # neighbors for each node should scale with growth in overlay participation (e.g., should not be $O(N)$)
- ❑ DHT mechanism should be fully distributed (no centralized point that bottlenecks throughput or can act as single point of failure)
- ❑ DHT mechanism should gracefully handle nodes joining/leaving the overlay
 - need to repartition the range space over existing nodes
 - need to reorganize neighbor set
 - need bootstrap mechanism to connect new nodes into the existing DHT infrastructure

DHT API

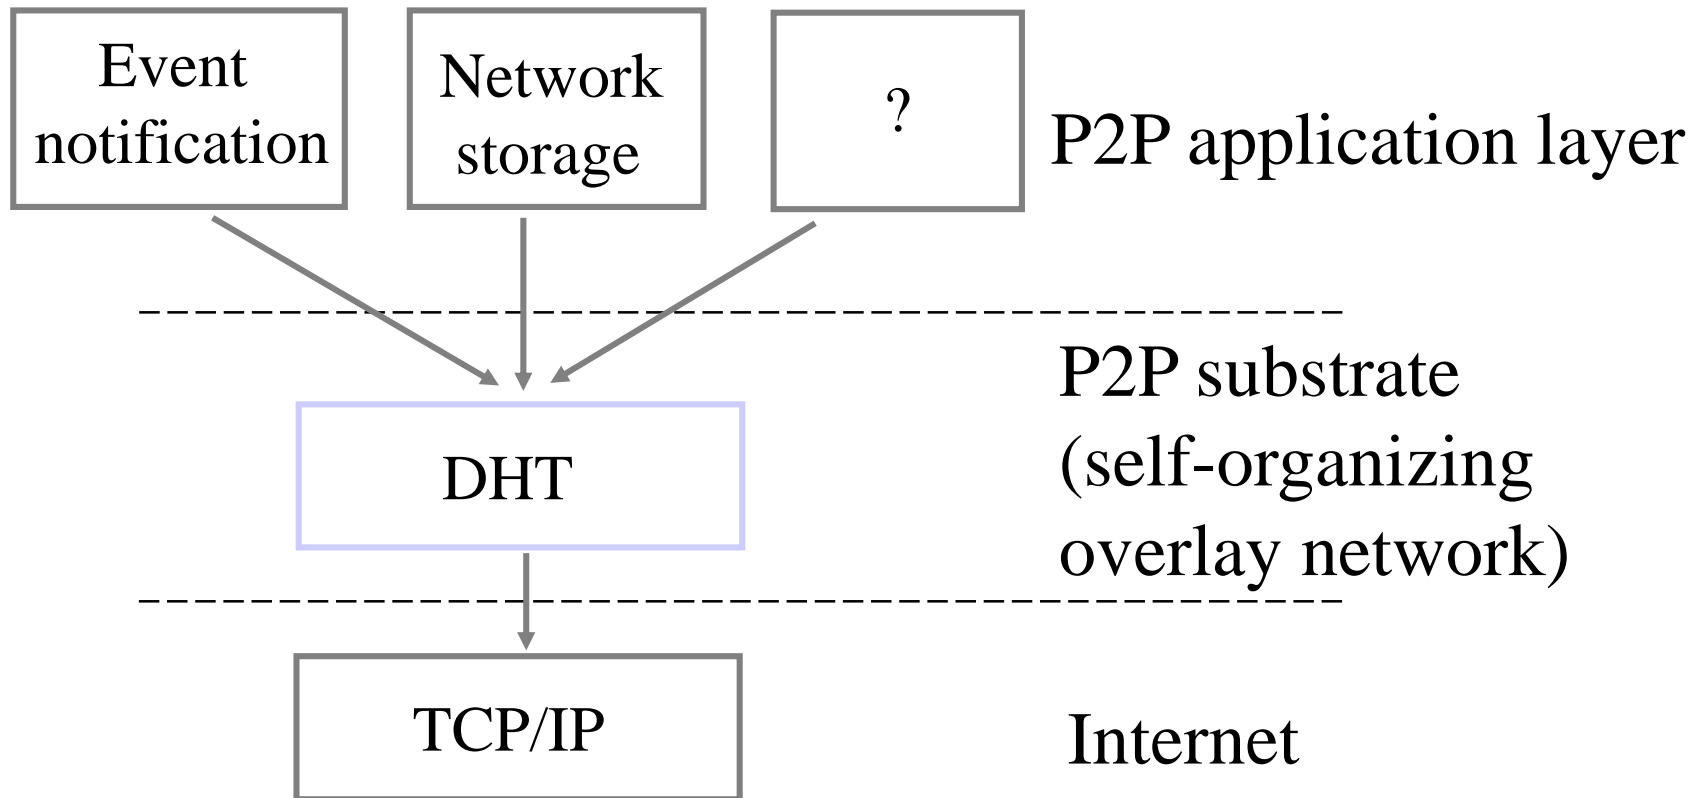
- ❑ each data item (e.g., file or metadata containing pointers) has a key in some ID space
- ❑ In each node, DHT software provides API:
 - Application gives API key k
 - API returns IP address of node that is responsible for k
- ❑ API is implemented with an underlying DHT overlay and distributed algorithms

DHT API

each data item (e.g., file or metadata pointing to file copies) has a key



DHT Layered Architecture



3. Structured P2P: DHT Approaches

- ❑ DHT service and issues
- ❑ **CARP**
- ❑ Consistent Hashing
- ❑ Chord
- ❑ CAN
- ❑ Pastry/Tapestry
- ❑ Hierarchical lookup services
- ❑ Topology-centric lookup service

CARP


DHT for cache clusters

- Each proxy has unique name

key = URL = u

- calc $h(\text{proxy}_n, u)$ for all proxies
- assign u to proxy with highest $h(\text{proxy}_n, u)$

if proxy added or removed, u is likely still in correct proxy



CARP (2)

- ❑ circa 1997
 - Internet draft:
Valloppillil and Ross
- ❑ Implemented in Microsoft & Netscape products
- ❑ Browsers obtain script for hashing from proxy automatic configuration file (loads automatically)

Not good for P2P:

- ❑ Each node needs to know name of all other up nodes
- ❑ i.e., need to know $O(N)$ neighbors
- ❑ But only $O(1)$ hops in lookup

3. Structured P2P: DHT Approaches


- ❑ DHT service and issues
- ❑ CARP
- ❑ Consistent Hashing
- ❑ Chord
- ❑ CAN
- ❑ Pastry/Tapestry
- ❑ Hierarchical lookup services
- ❑ Topology-centric lookup service

Consistent hashing (1)

- ❑ Overlay network is a circle
- ❑ Each node has randomly chosen id
 - Keys in same id space
- ❑ Node's successor in circle is node with next largest id
 - Each node knows IP address of its successor
- ❑ Key is stored in closest successor

Consistent hashing (2)

$O(N)$ messages
on avg to resolve
query



Note: no locality
among neighbors

Consistent hashing (3)

Node departures

- ❑ Each node must track $s \geq 2$ successors
- ❑ If your successor leaves, take next one
- ❑ Ask your new successor for list of its successors; update your s successors

Node joins


- ❑ You're new, node id k
- ❑ ask any node n to find the node n' that is the successor for id k
- ❑ Get successor list from n'
- ❑ Tell your predecessors to update their successor lists
- ❑ Thus, each node must track its predecessor

Consistent hashing (4)

- ❑ Overlay is actually a circle with small chords for tracking predecessor and k successors
- ❑ # of neighbors = $s+1$: $O(1)$
 - The ids of your neighbors along with their IP addresses is your "routing table"
- ❑ average # of messages to find key is $O(N)$

Can we do better?

Circular DHT with Shortcuts




- ❑ Each peer keeps track of IP addresses of predecessor, successor, short cuts.
- ❑ Reduced from 6 to 2 messages.
- ❑ Possible to design shortcuts so $O(\log N)$ neighbors, $O(\log N)$ messages in query

3. Structured P2P: DHT Approaches

- ❑ DHT service and issues
- ❑ CARP
- ❑ Consistent Hashing
- ❑ Chord
- ❑ CAN
- ❑ Pastry/Tapestry
- ❑ Hierarchical lookup services
- ❑ Topology-centric lookup service

Chord

- ❑ Nodes assigned 1-dimensional IDs in hash space at random (e.g., hash on IP address)
- ❑ Consistent hashing: Range covered by node is from previous ID up to its own ID (modulo the ID space)




Chord Routing

- ❑ A node s 's i^{th} neighbor has the ID that is equal to $s+2^i$ or is the next largest ID (mod ID space), $i \geq 0$
- ❑ To reach the node handling ID t , send the message to neighbor $\# \log_2(t-s)$
- ❑ Requirement: each node s must know about the next node that exists clockwise on the Chord (0^{th} neighbor)
- ❑ Set of known neighbors called a **finger table**

Chord Routing (cont'd)

- A node s is node t 's neighbor if s is the closest node to $t+2^i \bmod H$ for some i . Thus,
 - each node has at most $\log_2 N$ neighbors
 - for any object, the node whose range contains the object is reachable from any node in no more than $\log_2 N$ overlay hops
(each step can always traverse at least half the distance to the ID)
- Given K objects, with high probability each node has at most $(1 + \log_2 N) K / N$ in its range
- When a new node joins or leaves the overlay, $O(K / N)$ objects move between nodes




i	Finger table for node 67
0	72
1	72
2	72
3	86
4	86
5	1
6	32

Closest node clockwise to $67+2^i \bmod 100$

Chord Node Insertion


- ❑ One protocol addition: each node knows its closest counter-clockwise neighbor
- ❑ A node selects its unique (pseudo-random) ID and uses a bootstrapping process to find some node in the Chord
- ❑ Using Chord, the node identifies its successor in the clockwise direction
- ❑ An newly inserted node's predecessor is its successor's former predecessor



Example: Insert 82

Chord Node Insertion (cont'd)

- First: set added node s 's fingers correctly
 - s 's predecessor t does the lookup for each distance of 2^i from s



Lookups from node 72

Lookup(83) = 86 →

Lookup(84) = 86 →

Lookup(86) = 86 →

Lookup(90) = 1 →

Lookup(98) = 1 →


Lookup(14) = 32 →

Lookup(46) = 67 →

i	Finger table for node 82
0	86
1	86
2	86
3	1
4	1
5	32
6	67

Chord Node Insertion (cont'd)


- Next, update other nodes' fingers about the entrance of s (when relevant). For each i :
 - Locate the closest node to s (counter-clockwise) whose 2^i -finger can point to s : largest possible is $s - 2^i$
 - Use Chord to go (clockwise) to largest node t before or at $s - 2^i$
 - route to $s - 2^i$, if arrived at a larger node, select its predecessor as t
 - If t 's 2^i -finger routes to a node larger than s
 - change t 's 2^i -finger to s
 - set $t = \text{predecessor of } t$ and repeat
 - Else $i++$, repeat from top
- $O(\log^2 N)$ time to find and update nodes



e.g., for $i=3$

Chord Node Deletion

- Similar process can perform deletion



e.g., for $i=3$

3. Structured P2P: DHT Approaches

- ❑ DHT service and issues
- ❑ CARP
- ❑ Consistent Hashing
- ❑ Chord
- ❑ **CAN**
- ❑ Pastry/Tapestry
- ❑ Hierarchical lookup services
- ❑ Topology-centric lookup service

CAN


- ❑ hash value is viewed as a point in a D-dimensional cartesian space
- ❑ each node responsible for a D-dimensional "cube" in the space
- ❑ nodes are neighbors if their cubes "touch" at more than just a point (more formally, nodes s & t are neighbors when
 - s contains some $[<n_1, n_2, \dots, n_i, \dots, n_j, \dots, n_D>, <n_1, n_2, \dots, m_i, \dots, n_j, \dots, n_D>]$
 - and t contains $[<n_1, n_2, \dots, n_i, \dots, n_j + \delta, \dots, n_D>, <n_1, n_2, \dots, m_i, \dots, n_j + \delta, \dots, n_D>])$

				2			
		3		1	6	5	
				4			
7						8	
	30						

- Example: $D=2$
- 1's neighbors: 2,3,4,6
- 6's neighbors: 1,2,4,5
- Squares "wrap around", e.g., 7 and 8 are neighbors
- expected # neighbors: $O(D)$

CAN routing

- To get to $\langle n_1, n_2, \dots, n_D \rangle$ from $\langle m_1, m_2, \dots, m_D \rangle$
 - choose a neighbor with smallest cartesian distance from $\langle m_1, m_2, \dots, m_D \rangle$ (e.g., measured from neighbor's center)




- e.g., region 1 needs to send to node covering X
- checks all neighbors, node 2 is closest
- forwards message to node 2
- Cartesian distance monotonically decreases with each transmission
- expected # overlay hops: $(DN^{1/D})/4$

CAN node insertion

❑ To join the CAN:



- find some node in the CAN (via bootstrap process)
- choose a point in the space uniformly at random
- using CAN, inform the node that currently covers the space
- that node splits its space in half
 - 1st split along 1st dimension
 - if last split along dimension $i < D$, next split along $i+1$ st dimension
 - e.g., for 2-d case, split on x-axis, then y-axis
- keeps half the space and gives other half to joining node



Observation: the likelihood of a rectangle being selected is proportional to its size, i.e., big rectangles chosen more frequently


CAN node removal

- ❑ Underlying cube structure should remain intact
 - i.e., if the spaces covered by s & t were not formed by splitting a cube, then they should not be merged together
- ❑ Sometimes, can simply collapse removed node's portion to form bigger rectangle
 - e.g., if 6 leaves, its portion goes back to 1
- ❑ Other times, requires juxtaposition of nodes' areas of coverage
 - e.g., if 3 leaves, should merge back into square formed by 2,4,5
 - cannot simply collapse 3's space into 4 and/or 5
 - one solution: 5's old space collapses into 2's space, 5 takes over 3's space



CAN (recovery from) removal process

7		4	2 5		12		11
8	9 10	3		1	6		
13					14		



- View partitioning as a binary tree of
 - leaves represent regions covered by overlay nodes (labeled by node that covers the region)
 - intermediate nodes represent "split" regions that could be "reformed", i.e., a leaf can appear at that position
 - siblings are regions that can be merged together (forming the region that is covered by their parent)

3. Structured P2P: DHT Approaches

- ❑ DHT service and issues
- ❑ CARP
- ❑ Consistent Hashing
- ❑ Chord
- ❑ CAN
- ❑ Pastry
- ❑ Hierarchical lookup services
- ❑ Topology-centric lookup service

Routing table (node: 65a1fc04)

Row 0

<i>0</i>	<i>1</i>	<i>2</i>	<i>3</i>	<i>4</i>	<i>5</i>		<i>7</i>	<i>8</i>	<i>9</i>	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>	<i>f</i>
<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>		<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>

Row 1

<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>		<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>
<i>0</i>	<i>1</i>	<i>2</i>	<i>3</i>	<i>4</i>		<i>6</i>	<i>7</i>	<i>8</i>	<i>9</i>	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>	<i>f</i>
<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>		<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>

Row 2

<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>		<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>
<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>		<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>
<i>0</i>	<i>1</i>	<i>2</i>	<i>3</i>	<i>4</i>	<i>5</i>	<i>6</i>	<i>7</i>	<i>8</i>	<i>9</i>		<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>	<i>f</i>
<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>		<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>

Row 3

<i>6</i>		<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>	<i>6</i>
<i>5</i>		<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>	<i>5</i>
<i>a</i>		<i>a</i>	<i>a</i>	<i>a</i>	<i>a</i>	<i>a</i>	<i>a</i>	<i>a</i>	<i>a</i>	<i>a</i>	<i>a</i>	<i>a</i>	<i>a</i>	<i>a</i>	<i>a</i>
<i>0</i>		<i>2</i>	<i>3</i>	<i>4</i>	<i>5</i>	<i>6</i>	<i>7</i>	<i>8</i>	<i>9</i>	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>	<i>f</i>
<i>x</i>		<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>	<i>x</i>

$\log_{16} N$
rows

Pastry: Routing procedure

```
if (destination is within range of our leaf set)
    forward to numerically closest member
else
    if (there's a longer prefix match in table)
        forward to node with longest match
    else
        forward to node in table
            (a) shares at least as long a prefix
            (b) is numerically closer than this node
```


3. Structured P2P: DHT Approaches

- ❑ The DHT service and API
- ❑ CARP
- ❑ Consistent Hashing
- ❑ Chord
- ❑ CAN
- ❑ Pastry/Tapestry
- ❑ Hierarchical lookup services


Hierarchical Lookup Service

- ❑ KaZaA is hierarchical but unstructured
- ❑ Routing in Internet is hierarchical
- ❑ Perhaps DHTs can benefit from hierarchies too?
 - Peers are organized into groups
 - Inter-group lookup, then intra-group lookup

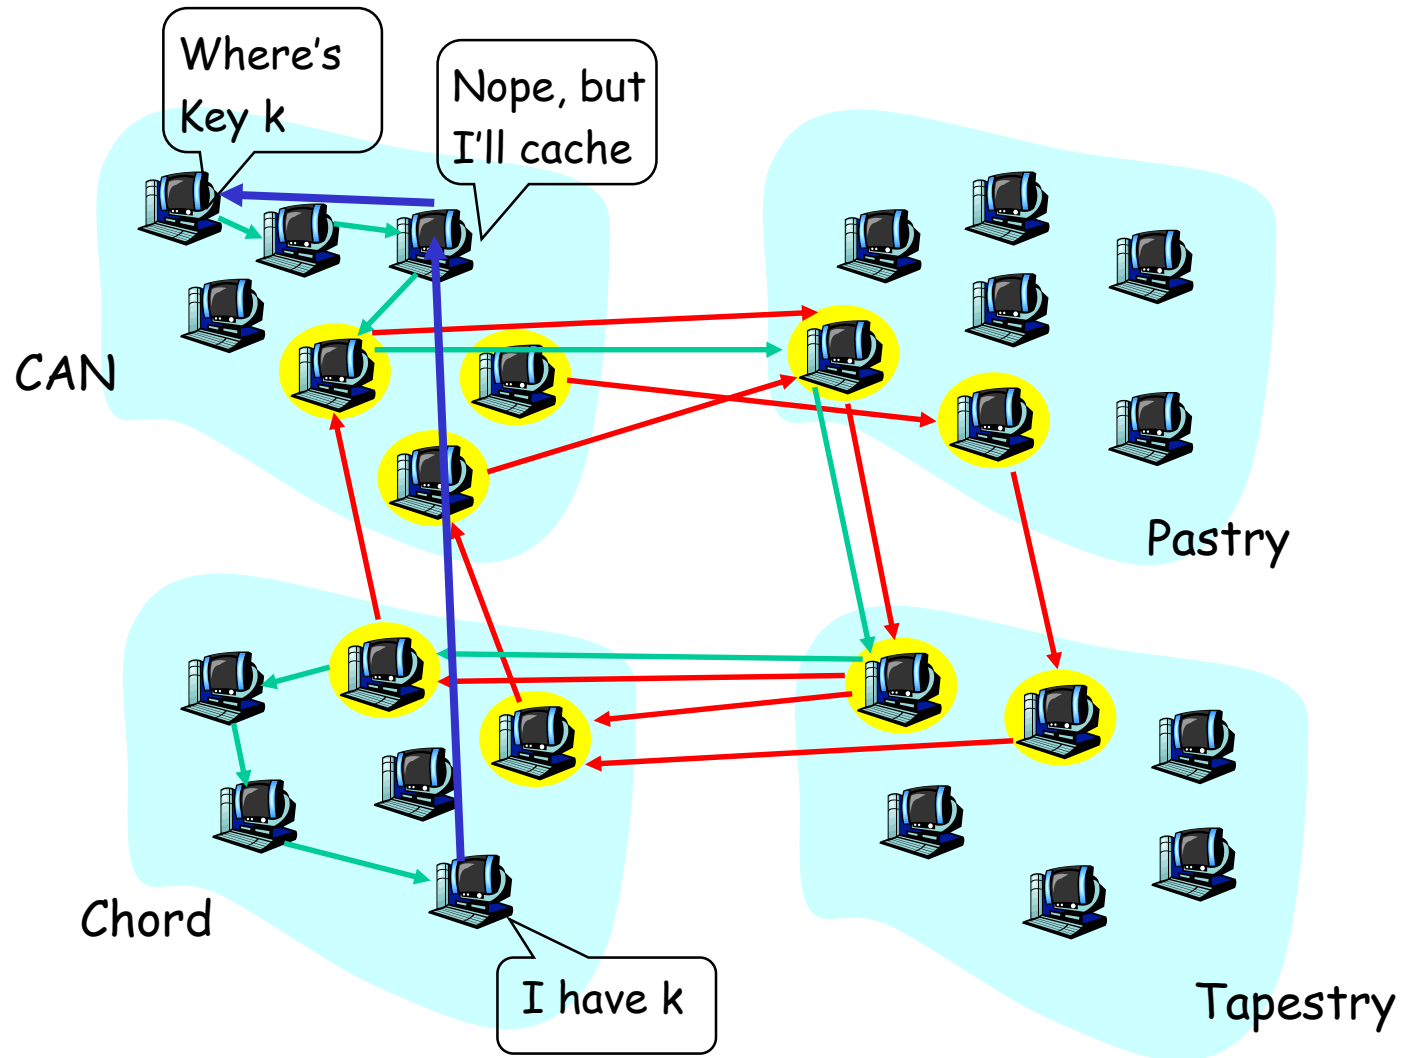
Hierarchical framework



Hierarchical Lookup



Cooperative group caching



Hierarchical Lookup (2)

Benefits


- ❑ Can reduce number of hops, particularly when nodes have heterogeneous availabilities
- ❑ Groups can cooperatively cache popular files, reducing average latency
- ❑ Facilitates large scale deployment by providing administrative autonomy to groups
 - Each ISP can use its own DHT protocol
 - Similar to intra-AS routing

SplitStream Multicast

- Balance load over peers
- Accommodate different limitations
 - Each node has a desired indegree and a forwarding capacity (max outdegree)
- Be robust to failures

The SplitStream approach

- Split data into stripes, each over its own tree
- Each node is internal to only one tree
- Built on Pastry and Scribe
 - Recall that Pastry uses prefix routing



Scribe background

- Built on top of Pastry
- Any Scribe node may create a group
 - Other nodes may join group or send multicast
- Node with nodeId numerically closest to groupId is the *rendezvous point*
 - Root of multicast tree for the group
 - Joins handled locally
- But it's only a single tree

Stripes

- SplitStream divides data into stripes
- Each stripe uses one Scribe multicast tree
- Prefix routing ensures property that each node is internal to only one tree
 - Inbound bandwidth: can achieve desired indegree while this property holds
 - Outbound bandwidth: this is harder—we'll have to look at the node join algorithm to see how this works

Respecting forwarding capacity

- The tree structure described may not respect maximum capacities
- Scribe's push-down fails to resolve the problem because a leaf node in one tree may have children in another tree